

## Problem Statement:

We are working with a car salesman to develop a model to predict the total dollar amount that customers are willing to pay given the following attributes:

1. Customer Name
2. Customer e-mail
3. Country
4. Gender
5. Age
6. Annual Salary
7. Credit Card Debt
8. Net Worth

The model will predict Car Purchase Amount

This is a regression task, we are predicting a continuous value.

In [1]: *# Step 1: Importing libraries and dataset*

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

In [2]: `df = pd.read_csv(r'C:\Users\jihad\Desktop\ML\CarSalePrediction\Car_Purchasing_Data.csv')`

In [3]: `df.head()`

Out[3]:

	Customer Name	Customer e-mail	Country	Gender	Age	Annual Salary
0	Martina Avila	cubilia.Curae.Phasellus@quisaccumsanconvallis.edu	Bulgaria	0	41.851720	62812.09301
1	Harlan Barnes	eu.dolor@diam.co.uk	Belize	0	40.870623	66646.89292
2	Naomi Rodriquez	vulputate.mauris.sagittis@ametconsectetueradip...	Algeria	1	43.152897	53798.55112
3	Jade Cunningham	malesuada@dignissim.com	Cook Islands	1	58.271369	79370.03798
4	Cedric Leach	felis.ullamcorper.viverra@egetmollislectus.net	Brazil	1	57.313749	59729.15130

In [4]: `df.describe()`

Out[4]:

	Gender	Age	Annual Salary	Credit Card Debt	Net Worth	Car Purchase Amount
<b>count</b>	500.000000	500.000000	500.000000	500.000000	500.000000	500.000000
<b>mean</b>	0.506000	46.241674	62127.239608	9607.645049	431475.713625	44209.799218
<b>std</b>	0.500465	7.978862	11703.378228	3489.187973	173536.756340	10773.178744
<b>min</b>	0.000000	20.000000	20000.000000	100.000000	20000.000000	9000.000000
<b>25%</b>	0.000000	40.949969	54391.977195	7397.515792	299824.195900	37629.896040
<b>50%</b>	1.000000	46.049901	62915.497035	9655.035568	426750.120650	43997.783390
<b>75%</b>	1.000000	51.612263	70117.862005	11798.867487	557324.478725	51254.709517
<b>max</b>	1.000000	70.000000	100000.000000	20000.000000	1000000.000000	80000.000000

In [5]:

```
# We have no missing values.

print(" \n Total NaN in each column in df : \n\n",
      df.isnull().sum())

print(" \n Total NaN values in df : \n\n",
      df.isnull().sum().sum())
```

Total NaN in each column in df :

```
Customer Name      0
Customer e-mail    0
Country            0
Gender             0
Age               0
Annual Salary      0
Credit Card Debt   0
Net Worth          0
Car Purchase Amount 0
dtype: int64
```

Total NaN values in df :

0

In [6]:

```
df.tail()
```

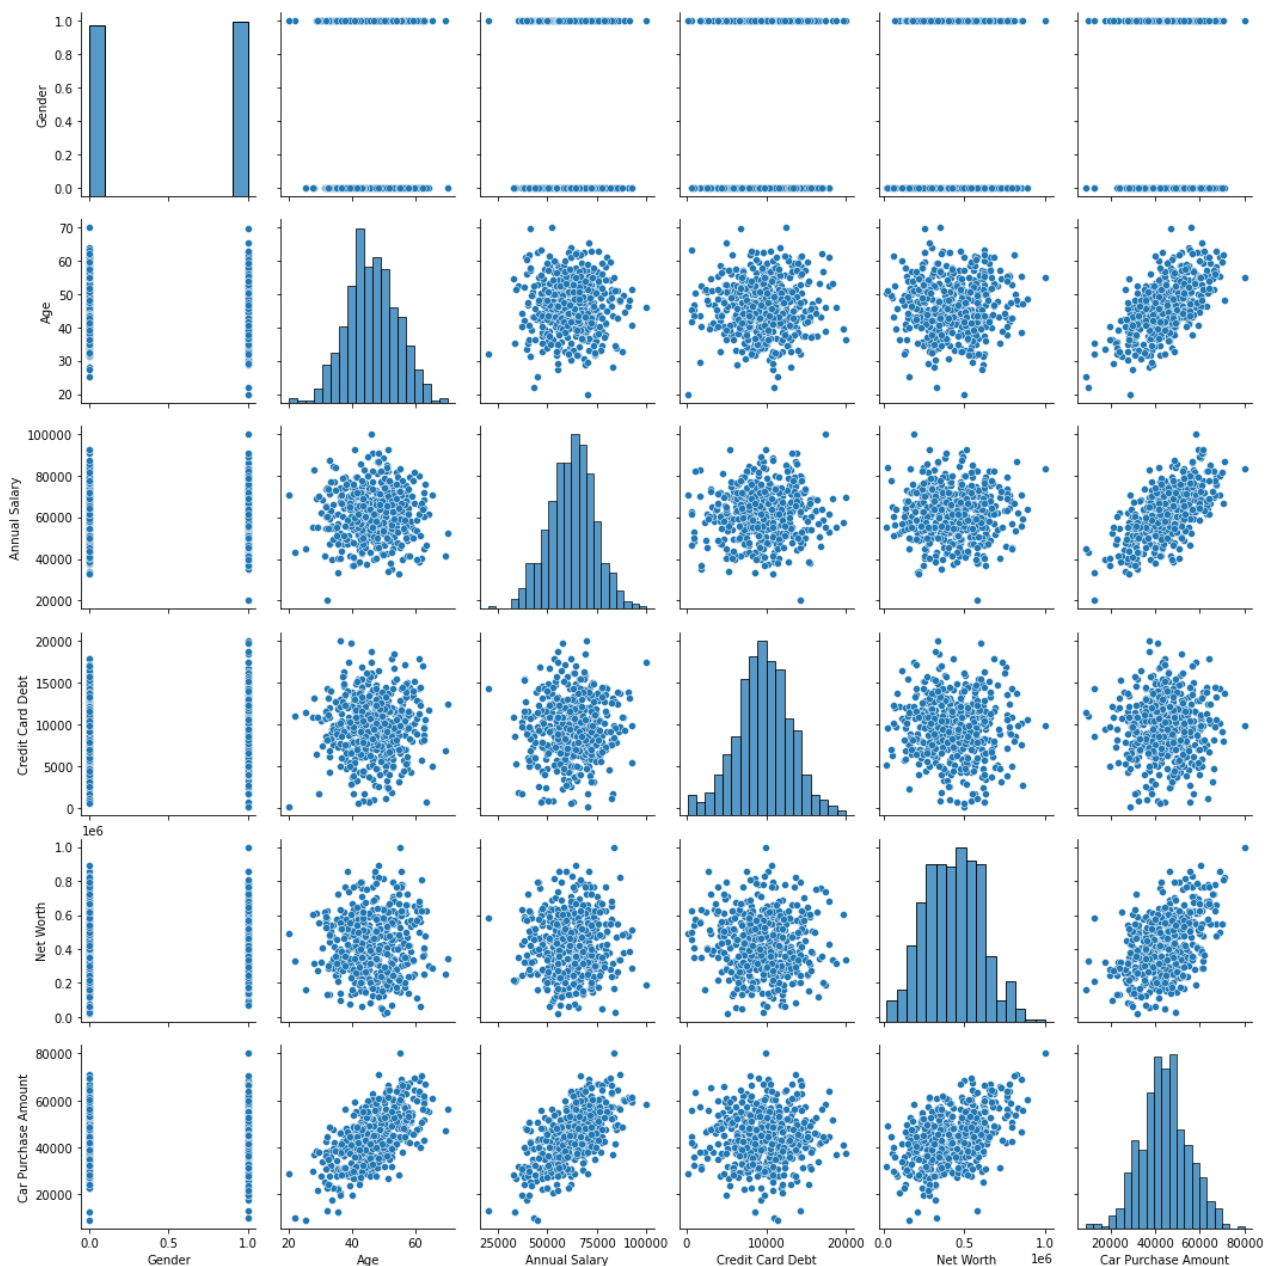
Out[6]:

	Customer Name	Customer e-mail	Country	Gender	Age	Annual Salary	Credit
<b>495</b>	Walter	ligula@Cumsociis.ca	Nepal	0	41.462515	71942.40291	6995.90
<b>496</b>	Vanna	Cum.sociis.natoque@Sedmolestie.edu	Zimbabwe	1	37.642000	56039.49793	12301.45
<b>497</b>	Pearl	penatibus.et@massanonante.com	Philippines	1	53.943497	68888.77805	10611.60
<b>498</b>	Nell	Quisque.varius@arcuVivamussit.net	Botswana	1	59.160509	49811.99062	14013.05
<b>499</b>	Marla	Camaron.marla@hotmail.com	marlal	1	46.731152	61370.67766	9391.34

```
In [7]: # Lets visualize the dataset.

# We can see correlation between age and purchase amount, annual salary and purchase am
# and net worth with purchase amount, which makes sense intuitively.
sns.pairplot(df)
```

```
Out[7]: <seaborn.axisgrid.PairGrid at 0x261f7014cc8>
```



```
In [8]: # Step 3: Create testing, training split and clean up the dataset.
# For our goal certain columns have no influence on the regression, so we're going to d
# Drop these columns (axis=1)
```

```
In [9]: # Step 3: Create testing, training split and clean up the dataset.
```

```
# the dataframe X will be the input variables from df we use in regression to predict Y
# Y= Car Purchase Amount.

# For our goal certain attributes are assumed to have no influence on the regression,
# we're going to drop these columns(axis=1) and the output ('Car Purchase Amount')

X = df.drop(['Customer Name', 'Customer e-mail', 'Country', 'Car Purchase Amount'], axis=1)
Y = df['Car Purchase Amount']
```

```
In [10]: Y.head() # Output we will train the model to predict
```

```
Out[10]: 0    35321.45877
         1    45115.52566
         2    42925.70921
         3    67422.36313
         4    55915.46248
         Name: Car Purchase Amount, dtype: float64
```

```
In [11]: X.head() # Inputs
```

```
Out[11]:
```

	Gender	Age	Annual Salary	Credit Card Debt	Net Worth
0	0	41.851720	62812.09301	11609.380910	238961.2505
1	0	40.870623	66646.89292	9572.957136	530973.9078
2	1	43.152897	53798.55112	11160.355060	638467.1773
3	1	58.271369	79370.03798	14426.164850	548599.0524
4	1	57.313749	59729.15130	5358.712177	560304.0671

```
In [12]: X.shape
```

```
Out[12]: (500, 5)
```

```
In [13]: Y.shape
```

```
Out[13]: (500,)
```

```
In [14]: # We now normalize the values

from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
X_scaled = scaler.fit_transform(X)
```

```
In [15]: # Maximum values

scaler.data_max_
```

```
Out[15]: array([1.e+00, 7.e+01, 1.e+05, 2.e+04, 1.e+06])
```

```
In [16]: # Minimum values
        scaler.data_min_
```

```
Out[16]: array([  0.,   20., 20000.,  100., 20000.])
```

```
In [17]: # In order for fit_transform to work with the output Y, we need to reshape Y (transpose)

        Y= Y.values.reshape(-1,1)
        Y.shape # shapes Y to (500,1)
```

```
Out[17]: (500, 1)
```

```
In [18]: Y_scaled = scaler.fit_transform(Y)
```

```
In [19]: # Step 4. Training the Model

        # We split the dataset into training and test datasets.
        # The test dataset is never seen by the model in the training.

        from sklearn.model_selection import train_test_split

        # 25% for testing data and 75% for the training data.
        X_train, X_test, y_train, y_test = train_test_split(X_scaled,Y_scaled,test_size=0.25)
```

```
In [20]: X_train.shape # 375 samples for testing
        X_test.shape  # 125 samples for testing
```

```
Out[20]: (125, 5)
```

```
In [21]: from tensorflow.keras import Sequential
        from tensorflow.keras.layers import Dense

        # All the outputs from one layer will be fully connected to the next (hidden) layer (dense)
        # input_dim = 5; 5 attribute columns we use for regression

        model = Sequential()

        # Specify inputs and neurons in the first layer, we also specify the activation function
        model.add(Dense(25, input_dim = 5, activation='relu'))

        # Lets add another hidden layer.
        # 25 neurons connected to the next 25 neurons
        model.add(Dense(25, activation='relu'))

        # We need to predict a value, we don't want to kill the output with impulse function
        model.add(Dense(1,activation='linear'))
```

```
In [22]: # Overview of parameters in model
        # We have a bias associated with each of the 25 neurons, so 150 = 125(X_samples) + 25(bias)
        # Next layer we have 25*25 = 625 neurons + 25(bias) = 650 param
```

```
# Output Layer 25 neurons connected to one output, (25 weights) + (1 bias) = 26 param
```

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 25)	150
dense_1 (Dense)	(None, 25)	650
dense_2 (Dense)	(None, 1)	26
Total params: 826		
Trainable params: 826		
Non-trainable params: 0		

In [23]:

```
# Training the model
```

```
# Using adam optimizier to speed up the gradient descent
```

```
# We take subsets of training data to avoid overfitting (validation_split)
```

```
# 100 epochs and a batch size of 50
```

```
model.compile(optimizer='adam', loss='mean_squared_error')
```

```
epochs_hist = model.fit(X_train, y_train, epochs = 100, batch_size = 50, verbose = 1, va
```

Epoch 1/100

6/6 [=====] - 0s 35ms/step - loss: 0.2100 - val\_loss: 0.1459

Epoch 2/100

6/6 [=====] - 0s 14ms/step - loss: 0.1230 - val\_loss: 0.0798

Epoch 3/100

6/6 [=====] - 0s 14ms/step - loss: 0.0624 - val\_loss: 0.0420

Epoch 4/100

6/6 [=====] - 0s 14ms/step - loss: 0.0313 - val\_loss: 0.0304

Epoch 5/100

6/6 [=====] - 0s 13ms/step - loss: 0.0230 - val\_loss: 0.0320

Epoch 6/100

6/6 [=====] - 0s 13ms/step - loss: 0.0232 - val\_loss: 0.0306

Epoch 7/100

6/6 [=====] - 0s 15ms/step - loss: 0.0210 - val\_loss: 0.0250

Epoch 8/100

6/6 [=====] - 0s 15ms/step - loss: 0.0176 - val\_loss: 0.0213

Epoch 9/100

6/6 [=====] - 0s 15ms/step - loss: 0.0161 - val\_loss: 0.0194

Epoch 10/100

6/6 [=====] - 0s 12ms/step - loss: 0.0154 - val\_loss: 0.0183

Epoch 11/100

6/6 [=====] - 0s 4ms/step - loss: 0.0149 - val\_loss: 0.0171

Epoch 12/100

6/6 [=====] - 0s 4ms/step - loss: 0.0141 - val\_loss: 0.0162

Epoch 13/100

6/6 [=====] - 0s 3ms/step - loss: 0.0135 - val\_loss: 0.0153

Epoch 14/100

6/6 [=====] - 0s 4ms/step - loss: 0.0129 - val\_loss: 0.0145

Epoch 15/100

6/6 [=====] - 0s 4ms/step - loss: 0.0124 - val\_loss: 0.0137

```
Epoch 16/100
6/6 [=====] - 0s 4ms/step - loss: 0.0118 - val_loss: 0.0130
Epoch 17/100
6/6 [=====] - 0s 4ms/step - loss: 0.0113 - val_loss: 0.0123
Epoch 18/100
6/6 [=====] - 0s 4ms/step - loss: 0.0107 - val_loss: 0.0117
Epoch 19/100
6/6 [=====] - 0s 4ms/step - loss: 0.0100 - val_loss: 0.0111
Epoch 20/100
6/6 [=====] - 0s 4ms/step - loss: 0.0093 - val_loss: 0.0105
Epoch 21/100
6/6 [=====] - 0s 4ms/step - loss: 0.0087 - val_loss: 0.0099
Epoch 22/100
6/6 [=====] - 0s 4ms/step - loss: 0.0081 - val_loss: 0.0091
Epoch 23/100
6/6 [=====] - ETA: 0s - loss: 0.008 - 0s 3ms/step - loss: 0.007
6 - val_loss: 0.0083
Epoch 24/100
6/6 [=====] - 0s 4ms/step - loss: 0.0071 - val_loss: 0.0076
Epoch 25/100
6/6 [=====] - 0s 3ms/step - loss: 0.0067 - val_loss: 0.0070
Epoch 26/100
6/6 [=====] - 0s 4ms/step - loss: 0.0062 - val_loss: 0.0064
Epoch 27/100
6/6 [=====] - 0s 4ms/step - loss: 0.0058 - val_loss: 0.0058
Epoch 28/100
6/6 [=====] - 0s 4ms/step - loss: 0.0054 - val_loss: 0.0053
Epoch 29/100
6/6 [=====] - 0s 4ms/step - loss: 0.0050 - val_loss: 0.0049
Epoch 30/100
6/6 [=====] - 0s 4ms/step - loss: 0.0047 - val_loss: 0.0046
Epoch 31/100
6/6 [=====] - 0s 4ms/step - loss: 0.0043 - val_loss: 0.0043
Epoch 32/100
6/6 [=====] - 0s 4ms/step - loss: 0.0039 - val_loss: 0.0039
Epoch 33/100
6/6 [=====] - 0s 4ms/step - loss: 0.0036 - val_loss: 0.0036
Epoch 34/100
6/6 [=====] - 0s 4ms/step - loss: 0.0034 - val_loss: 0.0034
Epoch 35/100
6/6 [=====] - 0s 4ms/step - loss: 0.0031 - val_loss: 0.0032
Epoch 36/100
6/6 [=====] - 0s 4ms/step - loss: 0.0029 - val_loss: 0.0030
Epoch 37/100
6/6 [=====] - 0s 4ms/step - loss: 0.0027 - val_loss: 0.0028
Epoch 38/100
6/6 [=====] - 0s 4ms/step - loss: 0.0025 - val_loss: 0.0028
Epoch 39/100
6/6 [=====] - 0s 4ms/step - loss: 0.0024 - val_loss: 0.0025
Epoch 40/100
6/6 [=====] - 0s 4ms/step - loss: 0.0021 - val_loss: 0.0024
Epoch 41/100
6/6 [=====] - 0s 3ms/step - loss: 0.0020 - val_loss: 0.0022
Epoch 42/100
6/6 [=====] - 0s 4ms/step - loss: 0.0018 - val_loss: 0.0021
Epoch 43/100
6/6 [=====] - 0s 4ms/step - loss: 0.0017 - val_loss: 0.0019
Epoch 44/100
6/6 [=====] - 0s 4ms/step - loss: 0.0016 - val_loss: 0.0018
Epoch 45/100
```

```
6/6 [=====] - 0s 4ms/step - loss: 0.0014 - val_loss: 0.0017
Epoch 46/100
6/6 [=====] - 0s 4ms/step - loss: 0.0013 - val_loss: 0.0015
Epoch 47/100
6/6 [=====] - 0s 4ms/step - loss: 0.0012 - val_loss: 0.0014
Epoch 48/100
6/6 [=====] - 0s 3ms/step - loss: 0.0011 - val_loss: 0.0012
Epoch 49/100
6/6 [=====] - 0s 5ms/step - loss: 0.0010 - val_loss: 0.0012
Epoch 50/100
6/6 [=====] - 0s 5ms/step - loss: 9.3841e-04 - val_loss: 0.0010
Epoch 51/100
6/6 [=====] - 0s 4ms/step - loss: 8.6705e-04 - val_loss: 9.5371
e-04
Epoch 52/100
6/6 [=====] - 0s 4ms/step - loss: 8.0008e-04 - val_loss: 8.6270
e-04
Epoch 53/100
6/6 [=====] - 0s 4ms/step - loss: 7.7435e-04 - val_loss: 7.9497
e-04
Epoch 54/100
6/6 [=====] - 0s 5ms/step - loss: 7.1351e-04 - val_loss: 7.7328
e-04
Epoch 55/100
6/6 [=====] - 0s 4ms/step - loss: 6.5005e-04 - val_loss: 6.8705
e-04
Epoch 56/100
6/6 [=====] - 0s 3ms/step - loss: 6.0174e-04 - val_loss: 6.6018
e-04
Epoch 57/100
6/6 [=====] - 0s 3ms/step - loss: 5.4491e-04 - val_loss: 5.8118
e-04
Epoch 58/100
6/6 [=====] - 0s 4ms/step - loss: 5.0000e-04 - val_loss: 5.6932
e-04
Epoch 59/100
6/6 [=====] - 0s 4ms/step - loss: 4.6675e-04 - val_loss: 5.1422
e-04
Epoch 60/100
6/6 [=====] - 0s 4ms/step - loss: 4.3191e-04 - val_loss: 4.7129
e-04
Epoch 61/100
6/6 [=====] - 0s 3ms/step - loss: 3.9898e-04 - val_loss: 4.5125
e-04
Epoch 62/100
6/6 [=====] - 0s 3ms/step - loss: 3.7095e-04 - val_loss: 4.0311
e-04
Epoch 63/100
6/6 [=====] - 0s 3ms/step - loss: 3.4847e-04 - val_loss: 3.7797
e-04
Epoch 64/100
6/6 [=====] - 0s 4ms/step - loss: 3.2403e-04 - val_loss: 3.6862
e-04
Epoch 65/100
6/6 [=====] - 0s 4ms/step - loss: 3.0569e-04 - val_loss: 3.4469
e-04
Epoch 66/100
6/6 [=====] - 0s 3ms/step - loss: 2.8711e-04 - val_loss: 3.1784
e-04
Epoch 67/100
```



```
6/6 [=====] - 0s 4ms/step - loss: 2.6911e-04 - val_loss: 3.1357e-04
Epoch 68/100
6/6 [=====] - 0s 4ms/step - loss: 2.5082e-04 - val_loss: 3.0055e-04
Epoch 69/100
6/6 [=====] - 0s 4ms/step - loss: 2.3513e-04 - val_loss: 2.8021e-04
Epoch 70/100
6/6 [=====] - 0s 4ms/step - loss: 2.2410e-04 - val_loss: 2.6324e-04
Epoch 71/100
6/6 [=====] - 0s 4ms/step - loss: 2.1308e-04 - val_loss: 2.5157e-04
Epoch 72/100
6/6 [=====] - 0s 4ms/step - loss: 2.0244e-04 - val_loss: 2.4065e-04
Epoch 73/100
6/6 [=====] - 0s 4ms/step - loss: 1.9522e-04 - val_loss: 2.3563e-04
Epoch 74/100
6/6 [=====] - 0s 4ms/step - loss: 1.8347e-04 - val_loss: 2.2579e-04
Epoch 75/100
6/6 [=====] - 0s 3ms/step - loss: 1.7611e-04 - val_loss: 2.1631e-04
Epoch 76/100
6/6 [=====] - 0s 4ms/step - loss: 1.7092e-04 - val_loss: 2.0513e-04
Epoch 77/100
6/6 [=====] - 0s 4ms/step - loss: 1.6995e-04 - val_loss: 2.0109e-04
Epoch 78/100
6/6 [=====] - 0s 4ms/step - loss: 1.6187e-04 - val_loss: 1.9605e-04
Epoch 79/100
6/6 [=====] - 0s 4ms/step - loss: 1.5434e-04 - val_loss: 1.8651e-04
Epoch 80/100
6/6 [=====] - 0s 3ms/step - loss: 1.4900e-04 - val_loss: 1.7947e-04
Epoch 81/100
6/6 [=====] - 0s 4ms/step - loss: 1.4076e-04 - val_loss: 1.8120e-04
Epoch 82/100
6/6 [=====] - 0s 3ms/step - loss: 1.4014e-04 - val_loss: 1.7969e-04
Epoch 83/100
6/6 [=====] - 0s 3ms/step - loss: 1.3537e-04 - val_loss: 1.7032e-04
Epoch 84/100
6/6 [=====] - 0s 4ms/step - loss: 1.2930e-04 - val_loss: 1.6850e-04
Epoch 85/100
6/6 [=====] - 0s 4ms/step - loss: 1.2765e-04 - val_loss: 1.6756e-04
Epoch 86/100
6/6 [=====] - 0s 3ms/step - loss: 1.2283e-04 - val_loss: 1.5755e-04
Epoch 87/100
```

```

6/6 [=====] - 0s 3ms/step - loss: 1.2125e-04 - val_loss: 1.5630
e-04
Epoch 88/100
6/6 [=====] - 0s 3ms/step - loss: 1.1873e-04 - val_loss: 1.5528
e-04
Epoch 89/100
6/6 [=====] - 0s 3ms/step - loss: 1.1594e-04 - val_loss: 1.5754
e-04
Epoch 90/100
6/6 [=====] - 0s 4ms/step - loss: 1.1236e-04 - val_loss: 1.5308
e-04
Epoch 91/100
6/6 [=====] - 0s 4ms/step - loss: 1.1158e-04 - val_loss: 1.4655
e-04
Epoch 92/100
6/6 [=====] - 0s 3ms/step - loss: 1.0745e-04 - val_loss: 1.4680
e-04
Epoch 93/100
6/6 [=====] - 0s 3ms/step - loss: 1.0725e-04 - val_loss: 1.4930
e-04
Epoch 94/100
6/6 [=====] - 0s 4ms/step - loss: 1.0550e-04 - val_loss: 1.4617
e-04
Epoch 95/100
6/6 [=====] - 0s 4ms/step - loss: 1.0351e-04 - val_loss: 1.4484
e-04
Epoch 96/100
6/6 [=====] - 0s 3ms/step - loss: 1.0160e-04 - val_loss: 1.3819
e-04
Epoch 97/100
6/6 [=====] - 0s 3ms/step - loss: 9.9596e-05 - val_loss: 1.3741
e-04
Epoch 98/100
6/6 [=====] - 0s 4ms/step - loss: 9.8326e-05 - val_loss: 1.4481
e-04
Epoch 99/100
6/6 [=====] - 0s 3ms/step - loss: 9.7393e-05 - val_loss: 1.3872
e-04
Epoch 100/100
6/6 [=====] - 0s 4ms/step - loss: 9.6212e-05 - val_loss: 1.3281
e-04

```

In [24]:

```

# Using matplotlib to visualize the epoch history


epochs_hist.history.keys()

plt.plot(epochs_hist.history['loss'])
plt.plot(epochs_hist.history['val_loss'])

plt.title('25 Neuron Model Loss Progress In Training')
plt.ylabel('Training and Validation Loss')
plt.xlabel('Epoch Number')
plt.legend(['Training Loss', 'Validation Loss'])

```

Out[24]: &lt;matplotlib.legend.Legend at 0x26181a76b88&gt;



In [25]:

```
# Changing the architecture of the network,
# Reducing the number of neurons in connected layers to 5

model5 = Sequential()
model5.add(Dense(5, input_dim = 5, activation='relu'))
model5.add(Dense(5, activation='relu'))

# Last layer
model5.add(Dense(1, activation='linear'))
model5.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
=====		
dense_3 (Dense)	(None, 5)	30
dense_4 (Dense)	(None, 5)	30
dense_5 (Dense)	(None, 1)	6
=====		
Total params: 66		
Trainable params: 66		
Non-trainable params: 0		

In [26]:

```
model5.compile(optimizer='adam', loss='mean_squared_error')
epochs_chart = model5.fit(X_train, y_train, epochs=100, batch_size=50, verbose=1,
```

```
Epoch 1/100
6/6 [=====] - 0s 16ms/step - loss: 0.0361 - val_loss: 0.0315
Epoch 2/100
6/6 [=====] - 0s 3ms/step - loss: 0.0262 - val_loss: 0.0264
Epoch 3/100
6/6 [=====] - 0s 3ms/step - loss: 0.0237 - val_loss: 0.0247
Epoch 4/100
6/6 [=====] - 0s 4ms/step - loss: 0.0219 - val_loss: 0.0231
Epoch 5/100
6/6 [=====] - 0s 3ms/step - loss: 0.0201 - val_loss: 0.0222
Epoch 6/100
```

```
6/6 [=====] - 0s 4ms/step - loss: 0.0189 - val_loss: 0.0218
Epoch 7/100
6/6 [=====] - 0s 4ms/step - loss: 0.0181 - val_loss: 0.0211
Epoch 8/100
6/6 [=====] - 0s 4ms/step - loss: 0.0174 - val_loss: 0.0203
Epoch 9/100
6/6 [=====] - 0s 3ms/step - loss: 0.0168 - val_loss: 0.0195
Epoch 10/100
6/6 [=====] - 0s 4ms/step - loss: 0.0160 - val_loss: 0.0187
Epoch 11/100
6/6 [=====] - 0s 4ms/step - loss: 0.0153 - val_loss: 0.0178
Epoch 12/100
6/6 [=====] - 0s 4ms/step - loss: 0.0144 - val_loss: 0.0169
Epoch 13/100
6/6 [=====] - 0s 3ms/step - loss: 0.0136 - val_loss: 0.0158
Epoch 14/100
6/6 [=====] - 0s 5ms/step - loss: 0.0127 - val_loss: 0.0147
Epoch 15/100
6/6 [=====] - 0s 3ms/step - loss: 0.0118 - val_loss: 0.0137
Epoch 16/100
6/6 [=====] - 0s 4ms/step - loss: 0.0110 - val_loss: 0.0127
Epoch 17/100
6/6 [=====] - 0s 4ms/step - loss: 0.0102 - val_loss: 0.0117
Epoch 18/100
6/6 [=====] - 0s 22ms/step - loss: 0.0095 - val_loss: 0.0109
Epoch 19/100
6/6 [=====] - 0s 4ms/step - loss: 0.0088 - val_loss: 0.0102
Epoch 20/100
6/6 [=====] - 0s 3ms/step - loss: 0.0082 - val_loss: 0.0095
Epoch 21/100
6/6 [=====] - 0s 4ms/step - loss: 0.0076 - val_loss: 0.0089
Epoch 22/100
6/6 [=====] - 0s 3ms/step - loss: 0.0071 - val_loss: 0.0083
Epoch 23/100
6/6 [=====] - 0s 4ms/step - loss: 0.0066 - val_loss: 0.0078
Epoch 24/100
6/6 [=====] - 0s 3ms/step - loss: 0.0062 - val_loss: 0.0073
Epoch 25/100
6/6 [=====] - 0s 4ms/step - loss: 0.0058 - val_loss: 0.0069
Epoch 26/100
6/6 [=====] - 0s 3ms/step - loss: 0.0054 - val_loss: 0.0065
Epoch 27/100
6/6 [=====] - 0s 4ms/step - loss: 0.0052 - val_loss: 0.0062
Epoch 28/100
6/6 [=====] - 0s 3ms/step - loss: 0.0049 - val_loss: 0.0058
Epoch 29/100
6/6 [=====] - 0s 4ms/step - loss: 0.0046 - val_loss: 0.0056
Epoch 30/100
6/6 [=====] - 0s 4ms/step - loss: 0.0044 - val_loss: 0.0054
Epoch 31/100
6/6 [=====] - 0s 4ms/step - loss: 0.0042 - val_loss: 0.0051
Epoch 32/100
6/6 [=====] - 0s 3ms/step - loss: 0.0040 - val_loss: 0.0049
Epoch 33/100
6/6 [=====] - 0s 4ms/step - loss: 0.0039 - val_loss: 0.0047
Epoch 34/100
6/6 [=====] - 0s 4ms/step - loss: 0.0037 - val_loss: 0.0045
Epoch 35/100
6/6 [=====] - 0s 4ms/step - loss: 0.0036 - val_loss: 0.0043
Epoch 36/100
```

```
6/6 [=====] - 0s 3ms/step - loss: 0.0034 - val_loss: 0.0042
Epoch 37/100
6/6 [=====] - 0s 4ms/step - loss: 0.0033 - val_loss: 0.0040
Epoch 38/100
6/6 [=====] - 0s 3ms/step - loss: 0.0032 - val_loss: 0.0039
Epoch 39/100
6/6 [=====] - 0s 4ms/step - loss: 0.0031 - val_loss: 0.0038
Epoch 40/100
6/6 [=====] - 0s 3ms/step - loss: 0.0030 - val_loss: 0.0036
Epoch 41/100
6/6 [=====] - 0s 3ms/step - loss: 0.0029 - val_loss: 0.0034
Epoch 42/100
6/6 [=====] - 0s 3ms/step - loss: 0.0028 - val_loss: 0.0034
Epoch 43/100
6/6 [=====] - 0s 3ms/step - loss: 0.0027 - val_loss: 0.0032
Epoch 44/100
6/6 [=====] - 0s 4ms/step - loss: 0.0026 - val_loss: 0.0031
Epoch 45/100
6/6 [=====] - 0s 3ms/step - loss: 0.0025 - val_loss: 0.0030
Epoch 46/100
6/6 [=====] - 0s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 47/100
6/6 [=====] - 0s 3ms/step - loss: 0.0024 - val_loss: 0.0028
Epoch 48/100
6/6 [=====] - 0s 4ms/step - loss: 0.0023 - val_loss: 0.0027
Epoch 49/100
6/6 [=====] - 0s 4ms/step - loss: 0.0022 - val_loss: 0.0026
Epoch 50/100
6/6 [=====] - 0s 4ms/step - loss: 0.0022 - val_loss: 0.0025
Epoch 51/100
6/6 [=====] - 0s 4ms/step - loss: 0.0021 - val_loss: 0.0024
Epoch 52/100
6/6 [=====] - 0s 3ms/step - loss: 0.0020 - val_loss: 0.0023
Epoch 53/100
6/6 [=====] - 0s 3ms/step - loss: 0.0020 - val_loss: 0.0023
Epoch 54/100
6/6 [=====] - 0s 3ms/step - loss: 0.0019 - val_loss: 0.0021
Epoch 55/100
6/6 [=====] - 0s 3ms/step - loss: 0.0019 - val_loss: 0.0021
Epoch 56/100
6/6 [=====] - 0s 3ms/step - loss: 0.0018 - val_loss: 0.0020
Epoch 57/100
6/6 [=====] - 0s 3ms/step - loss: 0.0017 - val_loss: 0.0019
Epoch 58/100
6/6 [=====] - 0s 4ms/step - loss: 0.0017 - val_loss: 0.0018
Epoch 59/100
6/6 [=====] - 0s 3ms/step - loss: 0.0016 - val_loss: 0.0018
Epoch 60/100
6/6 [=====] - 0s 4ms/step - loss: 0.0016 - val_loss: 0.0017
Epoch 61/100
6/6 [=====] - 0s 3ms/step - loss: 0.0015 - val_loss: 0.0016
Epoch 62/100
6/6 [=====] - 0s 3ms/step - loss: 0.0015 - val_loss: 0.0016
Epoch 63/100
6/6 [=====] - 0s 3ms/step - loss: 0.0014 - val_loss: 0.0015
Epoch 64/100
6/6 [=====] - 0s 4ms/step - loss: 0.0014 - val_loss: 0.0015
Epoch 65/100
6/6 [=====] - 0s 3ms/step - loss: 0.0013 - val_loss: 0.0014
Epoch 66/100
```

```
6/6 [=====] - 0s 4ms/step - loss: 0.0013 - val_loss: 0.0014
Epoch 67/100
6/6 [=====] - 0s 3ms/step - loss: 0.0012 - val_loss: 0.0013
Epoch 68/100
6/6 [=====] - 0s 3ms/step - loss: 0.0012 - val_loss: 0.0013
Epoch 69/100
6/6 [=====] - 0s 3ms/step - loss: 0.0012 - val_loss: 0.0012
Epoch 70/100
6/6 [=====] - 0s 3ms/step - loss: 0.0011 - val_loss: 0.0012
Epoch 71/100
6/6 [=====] - 0s 3ms/step - loss: 0.0011 - val_loss: 0.0011
Epoch 72/100
6/6 [=====] - 0s 3ms/step - loss: 0.0010 - val_loss: 0.0011
Epoch 73/100
6/6 [=====] - 0s 3ms/step - loss: 0.0010 - val_loss: 0.0010
Epoch 74/100
6/6 [=====] - 0s 3ms/step - loss: 9.7459e-04 - val_loss: 9.9572
e-04
Epoch 75/100
6/6 [=====] - 0s 4ms/step - loss: 9.4172e-04 - val_loss: 9.5484
e-04
Epoch 76/100
6/6 [=====] - 0s 3ms/step - loss: 9.0997e-04 - val_loss: 9.2351
e-04
Epoch 77/100
6/6 [=====] - 0s 4ms/step - loss: 8.8030e-04 - val_loss: 8.9230
e-04
Epoch 78/100
6/6 [=====] - 0s 3ms/step - loss: 8.5119e-04 - val_loss: 8.4638
e-04
Epoch 79/100
6/6 [=====] - 0s 3ms/step - loss: 8.2107e-04 - val_loss: 8.1767
e-04
Epoch 80/100
6/6 [=====] - 0s 3ms/step - loss: 7.9561e-04 - val_loss: 7.7663
e-04
Epoch 81/100
6/6 [=====] - 0s 3ms/step - loss: 7.6362e-04 - val_loss: 7.4292
e-04
Epoch 82/100
6/6 [=====] - 0s 3ms/step - loss: 7.3566e-04 - val_loss: 7.0983
e-04
Epoch 83/100
6/6 [=====] - 0s 3ms/step - loss: 7.1270e-04 - val_loss: 6.7627
e-04
Epoch 84/100
6/6 [=====] - 0s 4ms/step - loss: 6.8674e-04 - val_loss: 6.4332
e-04
Epoch 85/100
6/6 [=====] - 0s 4ms/step - loss: 6.6811e-04 - val_loss: 6.2181
e-04
Epoch 86/100
6/6 [=====] - 0s 3ms/step - loss: 6.4791e-04 - val_loss: 5.8497
e-04
Epoch 87/100
6/6 [=====] - 0s 3ms/step - loss: 6.2388e-04 - val_loss: 5.7475
e-04
Epoch 88/100
6/6 [=====] - 0s 3ms/step - loss: 5.9930e-04 - val_loss: 5.2880
e-04
```

```

Epoch 89/100
6/6 [=====] - 0s 3ms/step - loss: 5.7554e-04 - val_loss: 5.0272e-04
Epoch 90/100
6/6 [=====] - 0s 3ms/step - loss: 5.4818e-04 - val_loss: 4.9696e-04
Epoch 91/100
6/6 [=====] - 0s 4ms/step - loss: 5.3204e-04 - val_loss: 4.5658e-04
Epoch 92/100
6/6 [=====] - 0s 3ms/step - loss: 5.1086e-04 - val_loss: 4.3620e-04
Epoch 93/100
6/6 [=====] - 0s 4ms/step - loss: 4.9026e-04 - val_loss: 4.1858e-04
Epoch 94/100
6/6 [=====] - 0s 3ms/step - loss: 4.7570e-04 - val_loss: 3.9919e-04
Epoch 95/100
6/6 [=====] - 0s 4ms/step - loss: 4.5697e-04 - val_loss: 3.7777e-04
Epoch 96/100
6/6 [=====] - 0s 3ms/step - loss: 4.4237e-04 - val_loss: 3.5964e-04
Epoch 97/100
6/6 [=====] - 0s 4ms/step - loss: 4.2317e-04 - val_loss: 3.4093e-04
Epoch 98/100
6/6 [=====] - 0s 3ms/step - loss: 4.0967e-04 - val_loss: 3.2799e-04
Epoch 99/100
6/6 [=====] - 0s 3ms/step - loss: 3.9163e-04 - val_loss: 3.0729e-04
Epoch 100/100
6/6 [=====] - 0s 4ms/step - loss: 3.7913e-04 - val_loss: 2.9361e-04

```

```

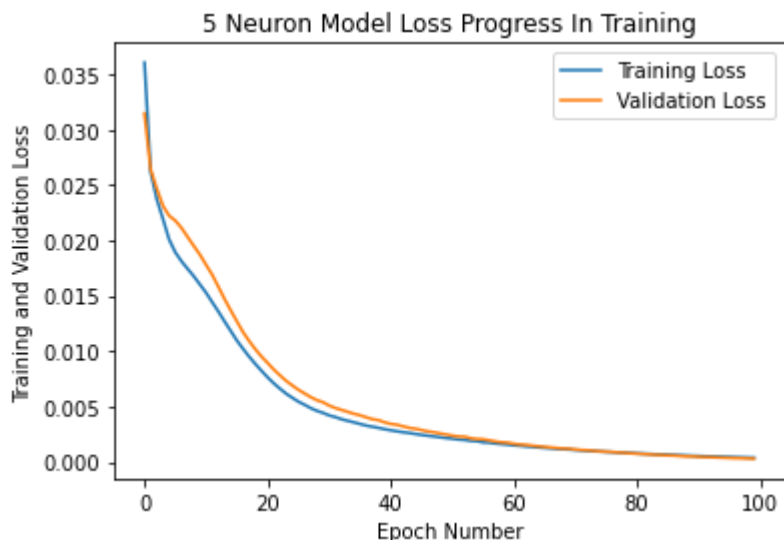
In [27]: plt.plot(epochs_chart.history['loss'])
plt.plot(epochs_chart.history['val_loss'])
plt.title('5 Neuron Model Loss Progress In Training')
plt.ylabel('Training and Validation Loss')
plt.xlabel('Epoch Number')
plt.legend(['Training Loss', 'Validation Loss'])

```

```

Out[27]: <matplotlib.legend.Legend at 0x2618177fd48>

```



In [28]:

```
# Lets predict with the models.
# We use scaled values for input and use an inverse transformation on the predicted val

# Test the model with the first row of X_scaled

# Gender, Age, Annual Salary, Credit Card Debt, Net Worth
X_test = np.array([[0, 0.4370344, 0.53515116, 0.57836085, 0.22342985]])

# Predictions
y_predict = model.predict(X_test)
y_predict5 = model5.predict(X_test)

# Printing values before inverting
print('(25NN) Expected Purchase Amount= ', y_predict)
print('(5NN) Expected Purchase Amount = ', y_predict5)

# Inverting predictions
y_predict_sample_actual = scaler.inverse_transform(y_predict)
y_predict_sample_actual5 = scaler.inverse_transform(y_predict5)
```

```
(25NN) Expected Purchase Amount= [[0.3685331]]
(5NN) Expected Purchase Amount = [[0.37696424]]
```

In [29]:

```
# Printing usable expected values from both models
print('(25NN) Actual Expected Purchase Amount= ', y_predict_sample_actual)
print('(5NN) Actual Expected Purchase Amount= ', y_predict_sample_actual5)
```

```
(25NN) Actual Expected Purchase Amount= [[35165.85]]
(5NN) Actual Expected Purchase Amount= [[35764.46]]
```

Conclusion:

For a customer with X\_test values for attributes, we show them cars in the price range of y\_predict.

In this way the model can help the salesmen specifically target customers better, leading to increased efficiency and more sales.



